

# iCaRL: Incremental Classifier and Representation Learning

## **Summary:**

This paper introduces iCaRL, a method that lets AI systems learn new object categories over time without forgetting the old ones—like how a child learns new animals without forgetting their pets. Traditional AI models need all data upfront and struggle to add new knowledge incrementally. iCaRL tackles this by combining three clever tricks:

- **Smart Memory (Exemplars):** Instead of storing every image it sees, iCaRL keeps a small set of key examples ("exemplars") for each class. These are chosen to best represent the class, like picking the most typical dog photos to remember what a dog looks like.
- **Adaptive Classification:** When classifying a new image, iCaRL compares it to the "average" of these exemplars (like matching a mystery animal to the average of stored dog/cat/bird examples). This avoids issues where updating the system messes up old knowledge.
- **Anti-Forgetting Training:** When learning new classes, iCaRL mixes new data with old exemplars and uses "knowledge distillation"—a technique where the system reviews its past decisions to avoid forgetting. Think of it as studying old notes while learning new material.

Tests on CIFAR-100 and ImageNet showed iCaRL outperforms other methods, especially when learning classes in small batches. Unlike rivals, it doesn't bias toward recent or early classes, balancing accuracy across all learned categories.

## **Implementation Details:**

**Network Architecture:** Uses a ResNet (a type of deep neural network) with 32 layers for CIFAR-100 and 18 layers for ImageNet. The network acts as a feature extractor, turning images into numerical vectors.

**Selection:** When a new class arrives, iCaRL picks exemplars iteratively. Each new exemplar is chosen to make the average of all exemplars match the average of the entire class. This is done via "herding," a method that prioritizes the most representative images.

**Storage:** Limits memory by capping exemplars (e.g., 2000 total for CIFAR-100). If space runs out, it trims older exemplars, keeping the most critical ones.

## **Training Summary:**

- **Mix Data:** Combine new class images with stored exemplars from old classes.
- **Avoid Forgetting:** Use "distillation loss" to make sure the network's predictions for old classes stay consistent with its past behavior.

- **Update Network:** Train the model on this mixed dataset, balancing new knowledge (classification loss) and old knowledge (distillation loss).
- **Classification:** For predictions, compute the average feature vector (prototype) of each class's exemplars. New images are labeled based on which prototype they're closest to.
- **Datasets:** Tested on CIFAR-100 (100 classes) and ImageNet (1000 classes), splitting classes into batches (e.g., 10 classes at a time).
- **Training:** Uses frameworks like TensorFlow/Theano. Learning rates start high (2.0) and drop over time. Training runs for 60–70 epochs per batch.
- **Results:** iCaRL beats alternatives like fine-tuning (which forgets old classes) or fixed-feature methods. For example, on CIFAR-100 with 10-class batches, iCaRL achieves ~64% accuracy vs. ~44% for the next-best method.

iCaRL mimics lifelong learning in humans, making AI systems more flexible and efficient. By smartly managing memory and balancing old/new knowledge, it avoids the "catastrophic forgetting" pitfall—a big step toward AI that learns continuously, like we do.

## Large Scale Incremental Learning

### **Summary:**

This paper tackles the challenge of catastrophic forgetting in AI models when learning new classes incrementally, especially on large datasets (like 1000+ classes). Traditional methods like iCaRL and EEIL work well on small datasets but fail at scale due to data imbalance (too many new vs. old samples) and visually similar classes (e.g., different dog breeds). The solution is A clever trick called Bias Correction (BiC).

**Problem:** The last layer of neural networks (FC layer) gets biased toward new classes because they have more training data.

**Solution:** Add a bias correction layer (just 2 parameters!) to adjust predictions for new classes. Train it on a small, balanced validation set split from old and new data.

Tested on ImageNet (1000 classes) and MS-Celeb-1M (10,000 faces) showed BiC outperforms iCaRL by ~26% and EEIL by ~18%. Even better, BiC's accuracy drop compared to a non-incremental model is just 10-16%, vs. 45% for older methods.

### **Implementation Details:**

#### Training Stage 1:

Train the model (ResNet) on new data + old exemplars.

Use distillation loss (to retain old knowledge) + classification loss (to learn new classes).

### Training Stage 2:

Freeze the model.

Add a linear bias correction layer after the FC layer.

Train this layer on a balanced validation set (split from exemplars and new data).

The layer adjusts new-class logits:

Copy

$$\text{adjusted\_score} = \alpha * \text{original\_score} + \beta$$

### Key Components:

Validation Set: Small subset (9:1 split of exemplars) ensures unbiased correction.

Bias Parameters ( $\alpha$ ,  $\beta$ ): Learned to reduce overconfidence in new classes.

### Model Setup:

Datasets: ImageNet-1000, MS-Celeb-1M (10k faces), CIFAR-100.

Architecture: ResNet-18 (ImageNet, Celeb) and ResNet-32 (CIFAR).

Exemplar Storage: Fixed memory (e.g., 20k for ImageNet).

Training: 100-250 epochs per incremental step, with learning rate decay.

### It solves:

Solves Imbalance: The validation set mirrors real-world class distribution.

Scalability: Just 2 parameters added, minimal computational overhead.

Robustness: Works even with visually similar classes (e.g., 10k faces in Celeb).

### Results:

ImageNet-1000: BiC achieves 73.2% accuracy vs. iCaRL's 44%.

Celeb-10000: BiC hits 88% vs. iCaRL's 65.5%.

Small Datasets: Matches iCaRL/EEIL on CIFAR-100 but shines at scale.

Removing bias correction drops accuracy by ~16%.

9:1 validation split works best; random vs. herding exemplar selection has minimal impact.

BiC makes incremental learning practical for real-world applications like face recognition (adding new users) or product categorization (new items), where datasets are massive and constantly growing. By fixing the FC layer's bias, it bridges the gap between incremental and non-incremental models, paving the way for lifelong learning AI.

## [Cost-Effective Active Learning for Deep Image Classification](#)

## **Summary:**

This paper tackles the challenge of training deep learning models with fewer labeled images by introducing CEAL (Cost-Effective Active Learning). Traditional methods require tons of labeled data, which is time-consuming and expensive. CEAL cleverly combines two strategies:

Smart Labeling: Instead of labeling everything, it asks humans to label only the most confusing images (like blurry faces or ambiguous objects).

Auto-Labeling: It automatically assigns labels to high-confidence images (clear, easy-to-recognize ones) using the model's predictions, saving manual effort.

By mixing human-labeled "hard" examples and auto-labeled "easy" examples, CEAL trains robust models with up to 36% fewer labels while matching or beating traditional methods. Tests on face recognition (CACD) and object categorization (Caltech-256) showed CEAL achieves high accuracy even with limited labeled data.

## **Implementation Details:**

Initializing: Start with a tiny labeled dataset (e.g., 10% of images). Train a basic CNN (like AlexNet for objects or a custom network for faces).

### **Sample Selection**

Uncertain Samples: Pick images where the model is least confident (e.g., predictions are close between two classes). These are sent for human labeling.

High-Confidence Samples: Auto-label images where the model is very sure (e.g., 95%+ confidence). These get pseudo-labels (e.g., "cat" or "dog") without human input.

Update Model: Fine-tune the CNN using both the newly labeled data and pseudo-labeled data. This helps the model learn better features and avoid overfitting.

Repeat: Gradually lower the confidence threshold for auto-labeling as the model improves, allowing more pseudo-labels over time.

## **Key Findings:**

Dynamic Threshold: Starts strict (only very confident images are auto-labeled) and relaxes as the model improves.

Error Control: Auto-labels have low error rates (under 5.5%), ensuring reliable training.

Flexible Architecture: Works with any CNN (tested on AlexNet and custom networks).

Training:

Used pre-trained AlexNet for Caltech-256, custom CNN for CACD.

Learning rates: 0.001 for most layers, 0.01 for the final layer.

Fine-tuned with batches of 2000–5000 images per iteration.

Results:

On CACD, CEAL achieved 91.5% accuracy with only 63% labeled data, vs. 81% for competitors.

On Caltech-256, it reached 73.8% accuracy with 78% labels, saving ~19% manual effort.

CEAL makes deep learning more practical by reducing reliance on expensive labeled data. It's like teaching a student to study smarter—focusing on tough problems while self-learning easier ones. This approach is especially useful for tasks where labeling is costly (e.g., medical images) or datasets keep growing.

## End-to-End Incremental Learning

### **Summary:**

This paper tackles the "catastrophic forgetting" problem in AI—where models forget old knowledge when learning new things—by introducing an end-to-end incremental learning method. Imagine teaching a robot to recognize dogs, then cats, without it forgetting dogs! The key ideas are:

**Smart Memory for Old Classes:** Stores a small set of the most representative old-class images (like keeping a photo album of past pets) using a technique called herding selection.

**Balanced Learning:** Mixes new data with old-class samples to avoid bias toward recent classes. Think of studying both new math topics and reviewing old ones regularly.

**Anti-Forgetting Training:** Uses two types of "learning reminders":

**Cross-entropy loss:** Teaches the model new classes.

**Distillation loss:** Preserves old knowledge by mimicking past predictions, like using old notes to retain memory.

Tests on CIFAR-100 and ImageNet showed this method outperforms others, especially when learning many small batches. For example, on CIFAR-100, it achieved ~64% accuracy vs. ~55% for older methods, even after learning 100 classes incrementally.

### **Implementation Details:**

#### Memory Management:

Herding Selection: When adding new classes, pick the most "average" images (e.g., typical dog photos) to store. Discards less representative ones to save space.

Fixed Memory Size: Caps stored old-class samples (e.g., 2000 for CIFAR-100). If full, removes the least critical ones.

#### Training Process:

Mix Data: Combine new-class images with stored old-class exemplars.

Dual Loss: Cross-entropy: Learn new classes ("This is a cat!"). Distillation: Preserve old knowledge ("Don't forget how dogs look!").

Update Everything: Unlike some methods, updates both the feature extractor (how images are processed) and classifier (decision-making) end-to-end.

Balanced Fine-Tuning: After training, re-train the model with a balanced subset (equal old/new samples) to fix class imbalance. Uses a smaller learning rate to avoid overwriting knowledge.

Dynamic Thresholds: Starts strict (only high-confidence images are auto-labeled) and relaxes as the model improves, allowing more pseudo-labels over time.

#### Model Setup:

Datasets: Tested on CIFAR-100 (100 object classes) and ImageNet (1000 classes), split into batches (e.g., 10 classes at a time).

Architecture: Uses ResNet (32-layer for CIFAR, 18-layer for ImageNet).

#### Results:

On ImageNet, achieved 69.4% accuracy with 100-class batches, beating iCaRL by ~7%.

#### Key Findings:

Unlike earlier methods (e.g., iCaRL) that train features and classifiers separately, this method updates them together end-to-end, making the model more robust and accurate. Think of it as learning to cook while improving knife skills simultaneously, instead of mastering knives first and recipes later.

Reduced catastrophic forgetting significantly—accuracy dropped only ~5% after 10 incremental steps vs. ~20% for older methods.

This approach lets AI systems learn continuously, like humans, without forgetting old tasks. It's efficient (uses limited memory) and flexible (works with any neural network).

Applications include robotics, personalized AI assistants, or systems needing frequent updates (e.g., new product recognition in retail).